

StereoQueerEval - Guidelines for Participants

Version 1
18 September 2026

Contents

1	Task Description	2
2	Dataset	2
2.1	Training and Test Sets	3
3	Labels	3
3.1	Subtask A – Stereotype Detection	3
3.2	Subtask B – Hate Speech Classification	3
3.3	Subtask C – Target and Identity Identification	4
4	Data Format	4
5	Submission Format	5
5.1	Submission for Subtask A	5
5.2	Submission for Subtask B	5
5.3	Submission for Subtask C	5
5.4	General Submission Requirements	6
6	How to Submit Your Runs	6
7	Evaluation	6
7.1	Subtask A	6
7.2	Subtask B	6
7.3	Subtask C	6
8	Baselines	7
8.1	Random Baseline	7
8.2	Majority-Class Baseline	7
8.3	Zero-Shot LLM Baseline	7
9	External Resources and Additional Data	8
10	Ethical Considerations	8

1 Task Description

The **StereoQueerEval** shared task focuses on the identification of stereotypes and hate speech targeting people belonging to the LGBTQIA+ community in online discourse. The task is based on comments collected from YouTube. Importantly, the comment should be interpreted together with the corresponding **video title and description**, which provide contextual information that may be necessary to determine whether a comment contains a stereotype or hateful content. The shared task consists of three independent subtasks:

- **Subtask A – Stereotype Detection.** Participants must determine whether a YouTube comment contains a stereotype targeting the LGBTQIA+ community. This is a binary classification task.
- **Subtask B – Hate Speech Classification.** Participants must classify the comment into one of three categories: explicit hate speech, implicit hate speech, or non-hateful/supportive content.
- **Subtask C – Target and Identity Identification.** For comments classified as hateful, participants must identify: (i) whether the target is an *individual* or a *group*, and (ii) which LGBTQIA+ identity or identities are targeted. Multiple identities may be associated with the same comment.

Participation in the three subtasks is independent: participants may submit systems to any subset of the subtasks. Please be noted that Subtask C is inherently linked to Subtask B, therefore, participation in both is highly encouraged.

2 Dataset

The StereoQueerEval dataset consists of YouTube comments collected through the official YouTube Data API v3. The data cover four languages:

- **English (EN)**
- **Italian (IT)**
- **Dutch (NL)**
- **Persian (FA)** (only test set - to be released in January)

For each language, comments were collected from selected news and media channels using a set of LGBTQIA+-related keywords. The keywords were translated and culturally adapted for each language. Each instance consists of three textual fields:

- **yt_title:** the title of the YouTube video;
- **yt_description:** the description of the video;
- **yt_comment:** the user-generated YouTube comment.

Participants may use all three fields as input to their systems. The training and test sets are constructed from different YouTube videos to avoid direct overlap between the two splits.

2.1 Training and Test Sets

For EN, IT, and NL, the final dataset contains:

- **EN**: train=2,989 and test=743;
- **IT**: train=2,400 and test=597;
- **NL**: train=2,238 and test=554;
- **FA**: train=0 and test=1,654.

The training data are released with gold annotations. Test will be released in January 2027, and labels will be withheld and will be used for the official evaluation. The exact number of instances may vary slightly across languages because of the data collection and annotation process. For Persian, we will release only the test set in January 2027, allowing for cross-lingual training due to the higher difficulty of collecting relevant data.

3 Labels

Each training instance is associated with one label for Subtask A, one label for Subtask B and, when applicable, a target annotation for Subtask C.

3.1 Subtask A – Stereotype Detection

The following two labels are used:

- **yes**: the comment contains a stereotype targeting LGBTQIA+ people;
- **no**: the comment does not contain such a stereotype.

A stereotype is understood as a generalized, oversimplified, or fixed belief about a group of people. In this task, the focus is on stereotypes targeting people belonging to the LGBTQIA+ community.

3.2 Subtask B – Hate Speech Classification

Each comment belongs to one of the following three categories:

- **yes_explicit**: explicit hateful content, including direct slurs, insults, or other overtly hateful expressions targeting LGBTQIA+ individuals or groups;
- **yes_implicit**: hateful content expressed indirectly, implicitly, or through coded language, including implied insults, derogatory stereotypes, and other forms of indirect hateful expression;
- **no**: neutral, supportive, or otherwise non-hateful content.

The distinction between explicit and implicit hate speech is based on the annotation guidelines provided to the human annotators.

3.3 Subtask C – Target and Identity Identification

Subtask C is evaluated only on comments annotated as hateful in the gold data. For each hateful comment, participants must predict two components:

1. the **target type**: individual or group;
2. one or more targeted **LGBTQIA+ identities**.

Code	Identity
l	Lesbian
g	Gay
b	Bisexual
t	Transgender
q	Queer
i	Intersex
a	Asexual/Aromantic
nb	Non-binary
lgbtqia+	LGBTQIA+ community as a whole

Table 1: Identity labels for Subtask C.

A comment may target more than one identity. For example, a hateful comment targeting gay and bisexual people may be annotated as:

`group_g,b`

Similarly, a comment targeting an individual lesbian person is represented as:

`individual_l`

The order of multiple identity labels does not affect evaluation. The available identity labels are reported in Table 1.

4 Data Format

The training data is provided in TSV format (UTF-8 encoding). Titles, descriptions and comments may contain newlines ("`\n`") inside quoted fields. We recommend reading the files with a proper TSV parser. Each row corresponds to one YouTube triplet and contains the contextual information and gold annotations. The exact column names and file format will be provided with the released dataset. A simplified example is shown in Table 2.

ID	Title	Description	Comment	Labels
train_IT_0006	Gay Pride a Napoli: "Un pride antifascista contro ogni forma d'odio"	A Napoli migliaia di persone hanno partecipato al Gay Pride. Tra i temi principali della manifestazione, oltre a quelli del movimento lgbtqia+, anche quelli dell'antifascismo e del sostegno alla popolazione palestinese.	Le lobby lgbt cariche di odio e le persone omosessuali sono cose ben diverse. Ancora con sto stereotipo sciocco degli omosessuali che sono tutti a sinistra? Ma aggiornatevi!!	yes, yes_explicit, group_lgbtqia+
train_EN_0001	GOP Congresswoman uses anti-trans slur	Rep. Mace challenged for anti-transgender slur during a hearing on USAID funding; she responded "I don't really care."	Reply defending Mace and praising her "leadership."	yes, yes_implicit, group_t

Table 2: Illustrative structure of the dataset.

The data will be made accessible to participants after they have read and agreed to the terms contained in the **Data Usage Agreement** form that can be found at this [link](#).

5 Submission Format

Participants must submit predictions for the test instances on the dedicated Codabench website (online in October 2026), where there will be a distinct submission for each subtask. The submission must contain the identifier of each test instance and the corresponding prediction. The test instance identifiers must be copied exactly from the provided test file. The text of the comment, video title, or video description must **not** be included in the prediction file.

5.1 Submission for Subtask A

For Subtask A, each line must contain:

```
StereoQueerEval_id    stereotype
```

where `stereotype` must be either `yes` or `no`.

For example:

```
EN_0001    yes
EN_0002    no
EN_0003    no
EN_0004    yes
```

5.2 Submission for Subtask B

For Subtask B, each line must contain:

```
StereoQueerEval_id    hate
```

where `hate` must be one of:

```
yes_explicit
yes_implicit
no
```

For example:

```
EN_0001    no
EN_0002    yes_explicit
EN_0003    yes_implicit
EN_0004    no
```

5.3 Submission for Subtask C

For Subtask C, predictions must contain the target type followed by the targeted identity code(s):

```
StereoQueerEval_id    target
```

For example:

```
EN_002    group_t
EN_007    individual_l
EN_014    group_g,b
EN_021    group_lgbtqia+
```

Multiple identities must be separated by commas. The order of the identity codes is not relevant. Subtask C predictions should only be provided for the hateful instances included in the Subtask C evaluation set.

5.4 General Submission Requirements

- Prediction files must contain one prediction for every required test instance.
- The `StereoQueerEval_id` must match the identifier in the test set.
- Do not modify or reorder the identifiers.
- Do not include the original text in the submission.
- Label strings must exactly match the labels defined in these guidelines.

6 How to Submit Your Runs

The official submission platform Codabench and submission procedure will be announced via the Google Group during the month of October 2026. Participants will be able to submit runs separately for each language and subtask. The maximum number of submissions per team and subtask will be communicated before the evaluation period.

7 Evaluation

We will produce a separate ranking for each subtask, with the possibility of filtering the ranking per language. All official evaluations will be performed automatically on Codabench (the platform will be up and running in October 2026). The primary evaluation metric for the shared task is the **macro-averaged F1-score**.

7.1 Subtask A

Subtask A is evaluated using macro-averaged F1 over the two stereotype classes: **yes**, **no**. The official score is:

$$\text{Macro-F1} = \frac{F1_{\text{yes}} + F1_{\text{no}}}{2}. \quad (1)$$

7.2 Subtask B

Subtask B is evaluated using macro-averaged F1 over the three hate speech classes: **non_hateful**, **hateful_implicit**, **hateful_explicit**. Thus, each of the three classes contributes equally to the final score, regardless of its frequency in the test set.

7.3 Subtask C

Subtask C contains two prediction dimensions: target type and target identity. The primary evaluation considers the **joint target label**, which combines the target type and all predicted identity labels. For example:

```
group_t
individual_l
group_g,b
```

The official Subtask C score is the macro-averaged F1 computed over these joint labels.

8 Baselines

To provide participants with reference points for system development, we provide three baselines:

1. a **random baseline**;
2. a **majority-class baseline**;
3. a **zero-shot LLM baseline** using `Qwen2.5-7B-Instruct`.

All three baselines are evaluated using the same labels and evaluation procedures as participant systems.

8.1 Random Baseline

The random baseline predicts uniformly from the valid label space. For Subtasks A and B, the model randomly selects one of the available classes. For Subtask C, the random baseline independently samples:

- a target type from `individual` and `group`;
- between one and three identity labels from the available identity vocabulary.

A fixed random seed is used to ensure reproducibility. The random baseline is intended only as a lower reference point and does not represent a competitive system.

8.2 Majority-Class Baseline

The majority-class baseline always predicts the most frequent label in the evaluation data. For Subtasks A and B, this corresponds to always predicting the most frequent stereotype or hate speech category. For Subtask C, the majority baseline predicts the most frequent complete target combination, i.e., the most frequent combination of target type and identity set. Majority-class frequencies are computed from the corresponding training data and then applied to the test set. This ensures that the baseline does not use test set gold labels.

8.3 Zero-Shot LLM Baseline

The third baseline uses `Qwen2.5-7B-Instruct` in a zero-shot setting. The model receives the task description and the relevant input text without being provided with task-specific demonstrations or supervised examples. For the zero-shot experiments:

- **Subtask A:** the model is asked to output either `yes` or `no`;
- **Subtask B:** the model is asked to output one of `yes_explicit`, `yes_implicit`, or `no`;
- **Subtask C:** the model is asked to output a target type (`individual` or `group`) and one or more identity names, which are subsequently mapped to the official identity codes.

Greedy decoding is used with `temperature=0` and `do_sample=False`. The same model is reused across all three subtasks. The LLM baseline is intended to provide a reproducible reference for zero-shot multilingual instruction-following performance.

9 External Resources and Additional Data

Participants are free to use external resources, including publicly available datasets, pretrained language models, multilingual models, lexicons, knowledge bases, and other resources, unless otherwise specified by the official SemEval rules. Participants should clearly document all external resources used in their system descriptions (February 2027) and final system papers (April 2027).

10 Ethical Considerations

The task concerns sensitive content involving sexual orientation, gender identity, stereotypes, discrimination, and hate speech. Participants should take appropriate care when developing and evaluating their systems. The dataset was collected using the official YouTube API. User-generated identifying information was anonymized during preprocessing. Participants should not attempt to identify individual users or recover removed personal information. The dataset may contain offensive, discriminatory, or otherwise distressing content. Participants are encouraged to consider these risks when inspecting the data and to adopt appropriate measures to protect the well-being of researchers working with the dataset.